

Where the revenue leaks, and how early you can see it

A B2B SaaS revenue engine built so it can be graded: five revenue leaks were planted in a simulated business and sealed away, and the warehouse, the monitor and the churn model had to find them blind.

Python, DuckDB, dbt, scikit-learn, XGBoost, Power BI · 1,107 accounts, Oct 2024 to Sep 2026 · reviewer's guide: [How it works](#)

5 of 5

planted leaks found by the monitor, blind

44%


of churns flagged by the model in the 60 days before

2.9x

churn rate in the riskiest 10% of accounts, out of time

13 of 13

injected data-problem types caught by tests



The Revenue page: \$1.72M MRR at Sep 2026 (\$20.62M ARR), 108% up on a year earlier, trailing-12-month NRR 85.1%.

The problem

Revenue in a subscription business rarely falls off a cliff; it leaks. A price change nudges small customers out at renewal. A broken release sends tickets up and usage down, and the cancellations arrive a month or two later. An account manager stops calling a book. A billing integration double-counts upgrades. A sales channel brings in customers who were never going to stay. Each one is visible in the data long before it is visible in the board deck, if the analytics are built to see it.

Most BI portfolio projects stop at showing MRR. The harder question is whether the analytics would actually *catch* a leak, and how fast. On real data that cannot be proved, because nobody knows what the true causes were. So this project starts from the other end: simulate a business with known leaks, seal the answer key, build the whole analytics stack without looking at it, then grade the result.

The business and the data

A day-by-day simulation of a B2B SaaS company over 24 months: 1,107 accounts across SMB, Mid-Market and Enterprise, three plans priced per seat, monthly and annual billing, eight account managers. Each account carries a hidden health value that drives its product use, support load, CSAT, expansion, downgrades and churn. The simulation writes the same exports a real company would have:

SOURCE	WHAT IT HOLDS	SIZE
Stripe webhooks and customers	subscription created, updated, deleted; invoices paid; MRR in cents	287 churns
Zendesk tickets	priority, category, first response, resolution, CSAT	17,176
Product telemetry	one row per user session with feature counts	1.9 million
CRM change log	account manager, CS tier and company-size changes with effective and recorded dates	SCD2 source

Five leaks were planted, each with a date, a scope and the signal it should leave. When an account churns, the simulation records which leak (if any) caused it. On top, 39,161 rows of real-world mess were injected across 13 kinds: duplicate events under new ids, missing account ids, timestamps in mixed time zones, out-of-range CSAT, tickets solved before they were opened, CRM changes recorded weeks late.

PLANTED LEAK	FROM	WHAT HAPPENS
INC-01 Starter price rise	Nov 2025	Starter goes from \$29 to \$39 per seat at each account's next renewal; some small customers leave
INC-02 Release 5.2	Mar 2026	Reports breaks for three weeks for accounts that rely on it: bug tickets, usage collapse, churn later
INC-03 Neglected book	Jun 2025	One Mid-Market account manager stops working their accounts; health erodes over months
INC-04 Duplicate webhooks	Jan 2026	A Stripe retry bug re-sends subscription changes under new event ids
INC-05 Weak partner cohort	Apr-Jun 2025	A partner launch brings in poorly qualified accounts whose use decays after onboarding

The pipeline

Raw files land as text in DuckDB. dbt types and cleans them in staging (flagging known mess with warn-level tests rather than hiding it), removes webhook duplicates and builds daily MRR state in an intermediate layer, then publishes a Kimball star schema: a daily MRR movement ledger, daily and weekly usage, support tickets with SLAs, an SCD Type 2 customer dimension and a fiscal calendar. Every core model has an enforced contract. A metrics layer adds the MRR waterfall, NRR and GRR, cohort retention, a weekly health score and the revenue leak monitor. A churn model reads the marts and writes its scores back; the Power BI project is generated from code and reads the same tables.


Only one script, `scorecard/score.py`, ever opens the answer key. Everything it grades ran without it.

Proving it

- **105 dbt nodes** pass on every build: models, unique and not-null keys, referential integrity to the dimensions, accepted values and ranges, SCD2 checks (no overlaps or gaps, exactly one current row), a ledger that must reconcile to subscription state, and a waterfall that must balance every month. 8 warn-level tests fire, one for each class of injected mess.
- **Mutation test, 14 of 14 caught.** The built warehouse is broken 14 different ways (a duplicate ledger row, overlapping SCD2 versions, a waterfall that no longer balances, a ticket pointing at a customer that does not exist...). Each break must fail a test, or the suite is not proving anything.
- **69 of 69 report numbers match.** The report is opened in its own Power BI window, every measure is evaluated in Power BI's engine, and its figures are compared with independent SQL, including the top-8 call list and the per-account tables.
- **Nothing truncated.** Every line of text in the report (titles, cards, table cells, slicer values, under every slicer combination) is measured with the real fonts against the space it has; tables must fit without scrollbars.
- **Byte-identical reruns.** Two complete runs from scratch produce the same 90 output files, raw data to Power BI project.

The revenue leak monitor

Six detectors run as one dbt model. Each tests a metric against its own trailing baseline using only data available at the time, and every alert carries the first date it could actually have fired (a monthly alert only once the month has closed).



Seed 42, detector v2: the first matching alert for each planted leak, in days after it started.

INCIDENT	LEAK	STARTED	DAYS	DETECTOR	WHAT THE MONITOR SAW
INC-01	Starter price rise	2025-11-01	29	churn_spike	24 of 212 starter customers churned vs 4.4 expected
INC-02	Release 5.2 breaks Reports	2026-03-10	5	feature_usage_drop	reports_run per active account fell to 42.0% of baseline
INC-03	Neglected account-manager book	2025-06-01	394	am_book_retention	Connie Lawrence (Mid-Market): 12.4% of book lost in 6 months vs 5.5% for peers
INC-04	Stripe re-sends webhooks	2026-01-15	3	duplicate_billing	3 duplicate subscription events removed
INC-05	Weak partner cohort	2025-04-01	306	weak_cohort	2025-04 cohort month-9 retention 63.0% vs 84.0% for earlier cohorts

3 alerts matched no planted leak (false alarms). 1 more matched none of the matching rules but, looking up the true causes behind them, were mostly incident-driven churn: real downstream signals, counted separately rather than as wins.

One revision, checked on unseen worlds

The first, blind version of the monitor used ad-hoc thresholds ("twice the average", "mean plus three standard deviations of eight weeks"). On seed 42 it found 4 of 5 leaks with 10 false alarms. Reading the false alarms showed the thresholds firing on small counts, and the cohort detector checking too few milestones to see the weak cohort. Version 2 replaced the thresholds with Poisson, binomial and two-proportion tests and checked cohorts at months 3, 6, 9 and 12. Because v2 was written after seeing seed 42, both versions were rerun on three simulated worlds nobody had looked at ([full notes](#)):

v1, ad-hoc thresholds

SEED	INC-01	INC-02	INC-03	INC-04	INC-05	FALSE ALARMS
42	29d	5d	394d	3d	missed	10
7	29d	5d	missed	3d	244d	10
2024	29d	5d	303d	3d	214d	14
31337	29d	5d	missed	3d	244d	13

v2, statistical tests

SEED	INC-01	INC-02	INC-03	INC-04	INC-05	FALSE ALARMS
42	29d	5d	394d	3d	306d	3
7	29d	5d	missed	3d	365d	1
2024	29d	5d	244d	3d	214d	6
31337	29d	5d	152d	3d	306d	8

Days from the start of each incident to its first matching alert. The price rise, the broken release and the duplicate webhooks are caught within a month in every world. The neglected book is the hardest: the account-manager test needs months of book history before it has the power to fire.

Data quality: what the tests caught


INJECTED PROBLEM	INJECTED	FLAGGED	HANDLED BY
Stripe webhooks re-sent under new ids (INC-04)	15	15	removed in int_subscription_events
Stripe events delivered twice, same id	51	51	deduplicated in staging
Plan ids with stray casing or spaces	288	288	lower/trim in staging
Stripe customers without an account id	20	20	matched on email domain
Tickets without an account	99	99	kept, mapped to unknown member -1
Ticket times in IST instead of UTC	380	380	converted to UTC
Tickets solved before they were opened	34	34	resolution nulled and flagged
CSAT scores outside 1-5	33	33	CSAT nulled and flagged
Ticket priorities with stray casing	846	846	lower/trim in staging
Sessions with no end time	18,806	18,774	duration nulled
Sessions ending before they start	3,703	3,703	duration nulled
Sessions logged twice	14,839	14,839	deduplicated in staging
CRM changes recorded 30+ days late	47	47	SCD2 uses effective date

Flagged can sit just below injected where two problems land on the same row (a session duplicated and also missing its end time is counted once).

The churn model

One row per paying account per week, built only from what was known that week: 30-day usage velocity, seat adoption, days since last session, tickets against the account's own baseline, high-priority and bug tickets, SLA breaches, CSAT level and trend, seat changes, a recent price rise, downgrades, days to renewal, and the account's segment, plan, channel and manager. The label: does the account cancel in the next 60 days. The model trains on weeks up to December 2025 and is tested on March to August 2026, out of time and straight through a broken release it never saw.

OUT-OF-TIME TEST, MAR-AUG 2026	LOGISTIC REGRESSION (CHOSEN)	XGBOOST
ROC AUC	0.73	0.71
Precision-recall AUC (base rate 6.6%)	0.19	0.14
Churn rate in the riskiest 10%	19.4% (2.9x)	
Share of all churners inside that 10%	29%	





Of 140 churns in the test period, the model had the account in its top 10% within the 60 days before for 44% of them (median warning 41 days); the rules-based health score flagged 11%.

The split by true cause is the useful part. Churn caused by a leak leaves a trail (falling usage, ticket spikes, a recent price rise) and the model picks most of it up. Organic churn, where a healthy-looking account simply leaves at renewal, is caught far less often (22%): there is little to see. The rules-based health score, with weights fixed by hand, catches much less than the model; it stays in the report because a CS team can read it, and the gap is shown rather than hidden. The strongest features are days to renewal, seat count, seat adoption, annual billing, monthly billing.


The report

Four pages, generated as a Power BI project (TMDL model, PBIR report) from `powerbi/build_pbip.py` : Revenue for the board view, Cohorts for retention and adoption, Accounts for the customer-success team, and Account to look at one customer in depth.





Accounts: 86 high-risk accounts hold \$192K of MRR, with \$151K expected to churn in 60 days. The call list ranks them by expected loss with the reason in plain words.



Account (type to search)

Woods, Calhoun and Schmidt

Woods, Calhoun and Schmidt · Enterprise, managed by Gina Moore

MRR, latest month

60-day churn probability

Health score, average

Tickets

Average CSAT

MRR lost

\$38,625

80%

79

180

3.82

\$0

Why the model flags it: 459 seats; Enterprise plan; Renewal in 59 days.

MRR at month end

Weekly sessions and active users

Health score by week (below 50 = At Risk)

Billing ledger: 6 latest MRR changes

Date	Movement	Reason	Plan	Seats	MRR before	MRR after
28 May 2026	Expansion	seat_change	enterprise	459	\$29,537	\$38,625
28 Oct 2025	Expansion	seat_change	enterprise	351	\$22,131	\$29,537
28 Aug 2025	Expansion	seat_change	enterprise	263	\$19,775	\$22,131
28 May 2025	Expansion	seat_change	enterprise	235	\$16,073	\$19,775
28 Nov 2024	New	new	enterprise	191	\$0	\$16,073

Support tickets: 6 newest

Opened	Priority	Category	Status	Hours to resolve	CSAT
28 Sep 2026	low	feature-request	solved	32.0	4
28 Sep 2026	low	integration	solved	33.2	
20 Sep 2026	low	integration	solved	68.1	4
18 Sep 2026	normal	integration	solved	6.4	
16 Sep 2026	normal	bug	solved	10.6	3
12 Sep 2026	normal	bug	solved	9.3	

Account: one customer's MRR history, usage, health trend, billing ledger and tickets, with the model's reasons.

Built as code

- One command, `python pipeline.py`, runs everything: simulation, dbt build and tests, model, exports, Power BI project, scorecard. CI runs it on every push, plus the mutation test.

- The Power BI report is generated: every page, visual and measure. Column types come from the same queries that write its data.
- This case study is rendered from the warehouse by `docs/render_docs.py` ; it fails if any figure is missing.

Limits and next steps

- The data is simulated. The scores measure how well the analytics find what was planted, not how a real business behaves. The answer key is only possible because of that.
- Incident dates are fixed across seeds; a stronger test would also randomise when and where the leaks happen.
- The account-manager detector finds a neglected book months late; a test that pools evidence over time (CUSUM) would likely find it sooner.
- The health score's weights are hand-set and weak. Calibrating it against the model's drivers, while keeping it readable, is the obvious next step.

Every figure here was computed from the warehouse by `docs/render_docs.py`. Code, data generator and report: see the repository README.